

TP1 - Optimizing Memory Access

Youness Anouar

January 29, 2026

1 Exercise 1: Impact of Compiler Optimization (-O0 vs -O2)

1.1 Experimental Results

1.1.1 -O0 Optimization Level

Stride	Sum	Time (ms)	Bandwidth (MB/s)
1	1e6	3.808	2003.5
2	1e6	4.162	1833.1
4	1e6	4.560	1673.1
8	1e6	8.314	917.7
16	1e6	26.637	286.4
20	1e6	30.002	254.3

Table 1: Memory access performance with -O0

1.1.2 -O2 Optimization Level

Stride	Sum	Time (ms)	Bandwidth (MB/s)
1	1e6	1.938	3936.7
2	1e6	1.619	4712.4
4	1e6	2.513	3036.0
8	1e6	5.653	1349.6
16	1e6	28.372	268.9
20	1e6	16.172	471.8

Table 2: Memory access performance with -O2

1.2 Performance Comparison




Figure 1: Execution time and bandwidth comparison between -O0 and -O2

1.3 Analysis

- **-O2 is significantly faster than -O0 for all strides.**
- For stride 1, execution time is reduced by nearly $2\times$, while bandwidth almost doubles.
- As stride increases, bandwidth drops for both versions due to reduced spatial locality.
- Despite this, **-O2 consistently achieves higher bandwidth and lower execution time.**

Conclusion: Compiler optimizations greatly improve loop execution and memory access efficiency, especially for small strides where cache locality is optimal.

2 Exercise 2: Matrix Multiplication Loop Ordering

2.1 Results

n	Time ijk (s)	Time ikj (s)	BW ijk (GB/s)	BW ikj (GB/s)
100	0.0102	0.0072	1.46	2.08
500	0.6616	0.4805	2.82	3.88
1000	6.4897	5.1358	2.30	2.90

Table 3: Matrix multiplication performance comparison




Figure 2: Performance comparison between ijk and ikj loop order

2.2 Explanation

- **ijk order:** accesses matrix B column-wise, causing poor cache locality.
- **ikj order:** inner loop iterates over contiguous memory, maximizing cache reuse.
- **ikj is consistently 30–50% faster** due to improved spatial and temporal locality.

Conclusion: Loop ordering has a major impact on performance; sequential memory access is critical for cache efficiency.

3 Exercise 3: Blocked Matrix Multiplication

3.1 Results

Block Size	Time (s)	MFLOPS	Bandwidth (MB/s)
1	39.07	55.0	0.64
8	3.28	655.4	7.68
16	3.00	715.5	8.39
32	3.42	628.5	7.37
64	4.63	463.7	5.43
1024	11.80	182.0	2.13

Table 4: Blocked matrix multiplication performance




Figure 3: Impact of block size on performance

3.2 Cache Analysis

The working set of the blocked algorithm is:

$$\text{Working set} = 3 \times b^2 \times 8 \text{ bytes}$$

For $b = 16$:

$$3 \times 16^2 \times 8 = 6144 \text{ bytes} \approx 6 \text{ KB}$$

This fits entirely within a typical 32 KB L1 cache, allowing:

- High temporal locality
- Reuse of matrices A and B from L1
- Minimal cache misses

Conclusion: Block size 16 is optimal because it maximizes cache utilization while avoiding cache overflow.

4 Exercise 4: Memory Leak Detection with Valgrind

4.1 Initial Run (Before Fix)

The program was analyzed using Valgrind with the following command:

```
valgrind --leak-check=full --track-origins=yes ./4
```

Valgrind reports a memory leak, as shown below:

```
==2167== HEAP SUMMARY:
==2167==     in use at exit: 20 bytes in 1 blocks
==2167==   total heap usage: 3 allocs, 2 frees, 1,064 bytes allocated
==2167==
==2167== 20 bytes in 1 blocks are definitely lost in loss record 1 of 1
==2167==    at 0x4848899: malloc
==2167==    by 0x109369: duplicate_array (4.c:33)
==2167==    by 0x109433: main (4.c:50)
==2167==
==2167== LEAK SUMMARY:
==2167==  definitely lost: 20 bytes in 1 blocks
==2167==  indirectly lost: 0 bytes in 0 blocks
==2167==  possibly lost: 0 bytes in 0 blocks
==2167==  still reachable: 0 bytes in 0 blocks
==2167== ERROR SUMMARY: 1 errors from 1 contexts
```

This leak is caused by dynamically allocated memory in `duplicate_array()` that was not released using `free()`.

4.2 After Fixing the Leak

After adding the missing `free(array_copy)` instruction, the program was re-analyzed with Valgrind:

```
valgrind --leak-check=full --track-origins=yes ./4
```

The corrected output is shown below:

```
==2175== HEAP SUMMARY:
==2175==     in use at exit: 0 bytes in 0 blocks
==2175==   total heap usage: 3 allocs, 3 frees, 1,064 bytes allocated
==2175==
==2175== All heap blocks were freed -- no leaks are possible
==2175==
==2175== ERROR SUMMARY: 0 errors from 0 contexts
```

4.3 Conclusion

The Valgrind analysis confirms that:

- The initial program contained a memory leak of 20 bytes.
- Proper deallocation using `free()` completely removed the leak.
- The final version is memory-safe with zero reported errors.

This demonstrates the importance of systematic memory management and the effectiveness of Valgrind for detecting memory leaks.

5 Exercise 5: HPL Benchmark

Given

The theoretical peak performance of a single CPU core is:

$$P_{\text{core}} = 70.4 \text{ GFLOP/s}$$

The measured performance is:

$$P_{\text{HPL}} = \text{GFLOPS reported by HPL}$$

The efficiency is defined as:

$$\eta = \frac{P_{\text{HPL}}}{P_{\text{core}}}$$

1. Comparison between measured performance and theoretical peak

For all experiments, the measured performance satisfies:

$$P_{\text{HPL}} < P_{\text{core}}$$

The measured performance ranges from approximately:

- **Minimum:** ~ 2 GFLOP/s for small values of N and NB
- **Maximum:** ~ 59 GFLOP/s for $N = 20000$ and $NB = 256$

This corresponds to about **3% to 84%** of the theoretical peak performance, depending on the values of N and NB .

2. Efficiency for each experiment

The efficiency is computed as:

$$\eta = \frac{P_{\text{HPL}}}{70.4}$$

3. Influence of N and NB

Effect of matrix size N

When the matrix size N increases:

- Execution time increases rapidly (approximately proportional to N^3)
- The achieved GFLOP/s increases
- The efficiency improves

Larger matrices reduce the relative impact of loop overhead, memory latency, and non-computational operations, allowing the processor to operate closer to its peak performance.

Effect of block size NB

Execution time

- Small NB leads to long execution times
- Increasing NB significantly reduces execution time
- For large NB , execution time stabilizes

N	NB	GFLOPS	Efficiency η
1000	1	4.1002	0.058
1000	2	7.4289	0.106
1000	4	11.971	0.170
1000	8	19.541	0.278
1000	16	28.732	0.408
1000	32	31.551	0.448
1000	64	31.347	0.445
1000	128	29.531	0.419
1000	256	24.276	0.345
5000	1	2.5190	0.036
5000	2	4.7826	0.068
5000	4	8.3574	0.119
5000	8	15.966	0.227
5000	16	29.124	0.414
5000	32	38.714	0.550
5000	64	45.684	0.649
5000	128	47.470	0.674
5000	256	47.338	0.672
10000	1	2.0441	0.029
10000	2	4.4606	0.063
10000	4	7.0258	0.100
10000	8	15.598	0.222
10000	16	26.956	0.383
10000	32	40.425	0.574
10000	64	47.655	0.677
10000	128	53.183	0.755
10000	256	53.284	0.757
20000	1	2.2882	0.033
20000	2	4.2097	0.060
20000	4	7.2870	0.104
20000	8	14.858	0.211
20000	16	27.143	0.385
20000	32	41.320	0.587
20000	64	50.863	0.723
20000	128	56.260	0.799
20000	256	59.036	0.838

Table 5: HPL performance and efficiency for all experiments

Performance (GFLOP/s)

- $NB = 1\text{--}4$: very low performance
- $NB = 8\text{--}32$: rapid increase in performance
- $NB \geq 64$: performance reaches a plateau

Larger block sizes improve cache reuse and allow BLAS kernels to better exploit vectorized instructions.

Block sizes giving best performance

N	Best NB (approx.)
1000	32–64
5000	128–256
10000	128–256
20000	256

The optimal block size increases with the matrix size N .

4. Why the measured performance is lower than the theoretical peak

The theoretical peak performance assumes ideal conditions that cannot be fully achieved in practice. The main limiting factors are:

- Memory hierarchy limitations, such as cache misses and limited memory bandwidth
- Non-computational instructions (loads, stores, branches)
- Algorithmic overhead in HPL, particularly during panel factorization
- Pipeline and scheduling inefficiencies due to instruction dependencies
- Operating system noise and background processes