

TP2

Youness Anouar

February 5, 2026

Contents

1	Exercise 1	2
1.1	Questions 1–2: Manual Unrolling Baseline	2
1.2	Question 3: Best Unroll Factor Stability	2
1.3	Questions 4–5: Manual Unrolling vs. -O2	2
1.4	Question 6: Effect of Data Type	2
1.5	Question 7: Theoretical Lower Bound (Bandwidth)	3
1.6	Question 8: Why Performance Improves Then Saturates	4
2	Exercise 2	4
2.1	Question 1: Performance at -O0 vs. -O2	4
2.2	Question 2: Analysis of Compiler Optimizations	4
2.3	Question 3: Manual Optimization vs. Compiler Optimization	5
3	Exercise 3	5
3.1	Question 1: Code Analysis	5
3.2	Question 2: Profiling with Callgrind	5
3.3	Question 3: Amdahl’s Law Prediction	5
3.4	Question 4: Consistency Across Problem Sizes	6
3.5	Question 5: Gustafson’s Law	6
4	Exercise 4	7
4.1	Question 1: Profiling Analysis (N=512)	7
4.2	Question 2: Amdahl’s Law – Strong Scaling (N=512)	8
4.3	Question 3: Gustafson’s Law – Weak Scaling (N=512)	9
4.4	Question 4: Consistency Analysis (N=128)	9
4.5	Comparison: Exercise 3 vs. Exercise 4	10

1 Exercise 1

1.1 Questions 1–2: Manual Unrolling Baseline

We sum an array of $N = 10^6$ doubles using a manual unroll factor $U \in [1, 32]$. The execution time $T(U)$ is shown in Fig. 1.

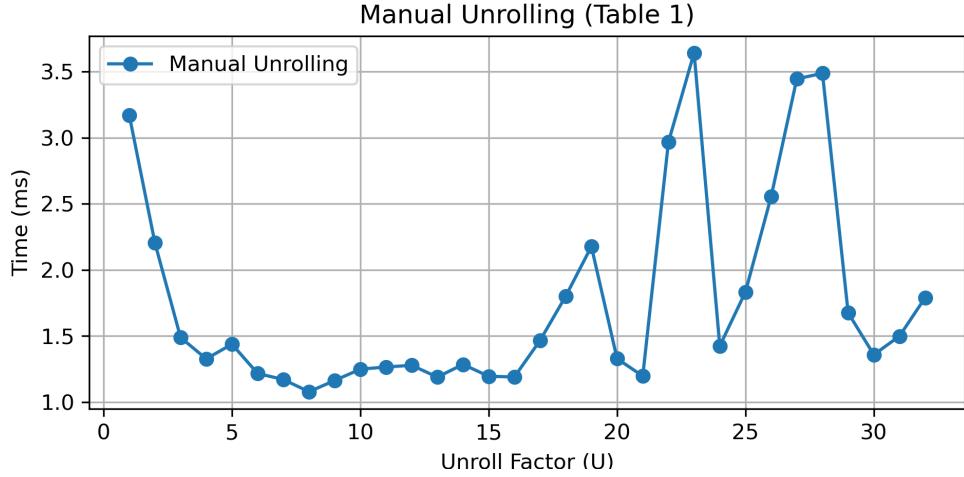


Figure 1: Manual unrolling with doubles

Observation:

- Time drops quickly when going from $U = 1$ to medium U .
- Best run in the shown experiment: around $U = 8\text{--}10$ (≈ 1.1 ms vs. 3.2 ms for $U = 1$).
- Very large U makes performance noisy and sometimes worse.

1.2 Question 3: Best Unroll Factor Stability

Across multiple runs the best U is not fixed; typical best values are in $[6, 10]$ (e.g., $U = 8$ in one run). Small timing noise and interactions with caches, branch prediction, and OS scheduling explain why the exact optimum moves, but the “good region” is stable.

1.3 Questions 4–5: Manual Unrolling vs. -O2

We now compare `-O0` and `-O2` while still manually unrolling. Results are summarized in Fig. 2.

Key points:

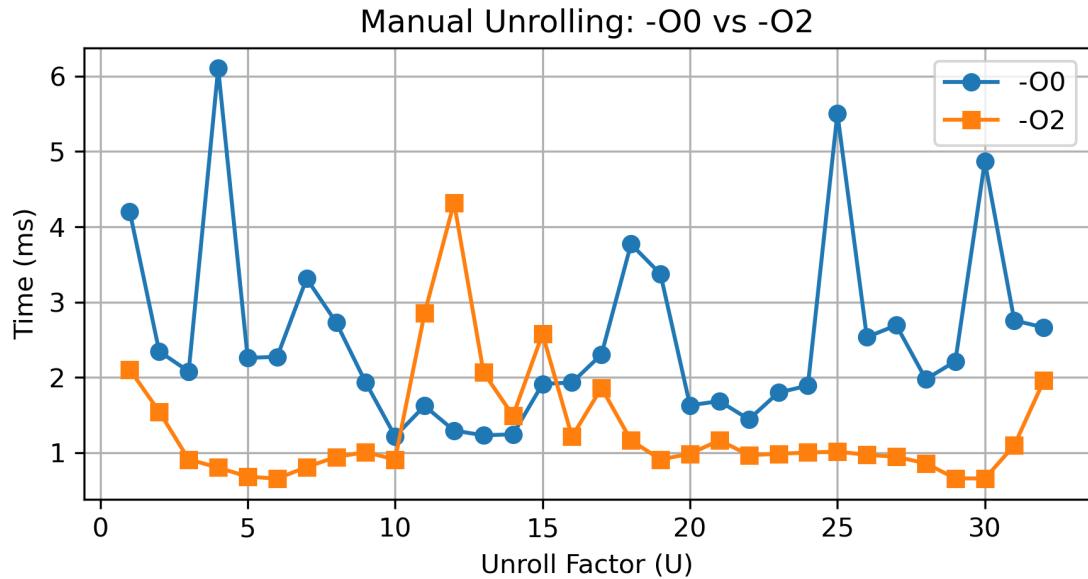
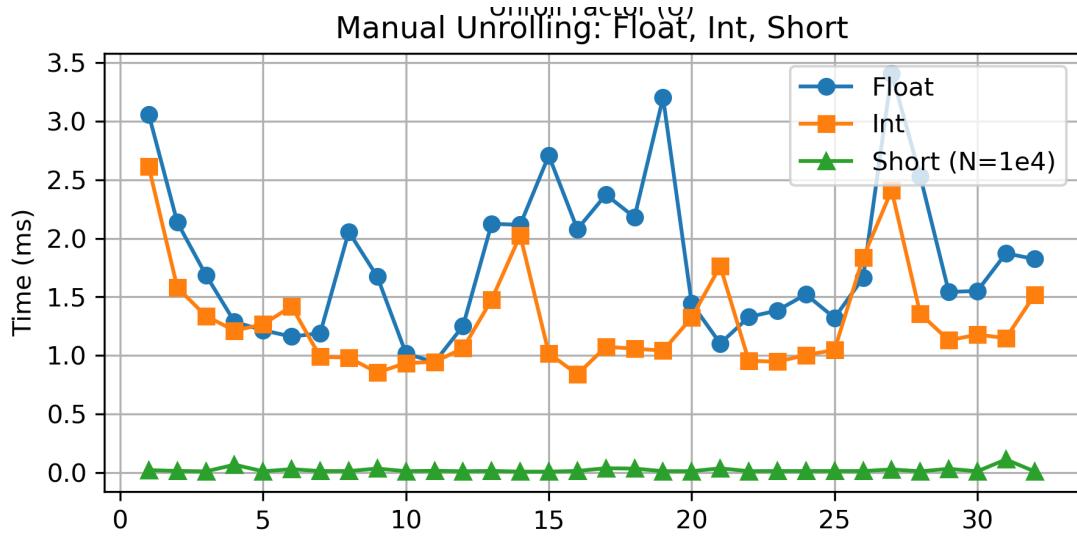
- `-O0`: best time ≈ 1.22 ms at $U = 10$ (about 65–70% faster than $U = 1$).
- `-O2`: best time ≈ 0.65 ms at $U = 6$ (also $\approx 70\%$ faster than $U = 1$ at `-O2`).
- Even with `-O2`, manual unrolling helps: the compiler does not fully remove loop overhead for this kernel.
- The best U with `-O2` is smaller (around 6) than with `-O0`, because the compiler already performs part of the work (strength reduction, instruction scheduling, partial unrolling).

Conclusion for Q5:

- Manual unrolling is clearly beneficial at `-O0`.
- With `-O2`, a moderate unroll factor plus compiler optimizations gives the best performance.
- The compiler alone does not completely replace manual unrolling here.

1.4 Question 6: Effect of Data Type

We repeat the experiment with `float`, `int`, and `short` (Fig. 3). For `short` we use $N = 10^4$ to avoid overflow.

**Figure 2:** Manual unrolling at -O0 vs. -O2.**Figure 3:** Manual unrolling for float, int, and short.

Observations:

- **Float:** best around $U = 11$ (≈ 0.94 ms).
- **Int:** best around $U = 16$ (≈ 0.84 ms).
- **Short ($N = 10^4$):** best around $U = 14\text{--}15$ (≈ 0.008 ms), but timings are very small and noisy.
- Optimal U depends on type; heavier types (double/float) hit compute and bandwidth limits earlier than int/short.

1.5 Question 7: Theoretical Lower Bound (Bandwidth)

We approximate a lower bound assuming the kernel is purely memory-bandwidth limited:

$$T_{\min} \approx \frac{N \times \text{sizeof(type)}}{\text{BW}},$$

with $N = 10^6$ (except short: $N = 10^4$) and $\text{BW} \approx 30 \text{ GB/s}$.

- **Double** (8 B): $8 \text{ MB} \Rightarrow T_{\min} \approx 0.27 \text{ ms}$.
Best measured $\approx 1.08 \text{ ms}$ ($\sim 4 \times$ slower).
- **Float** (4 B): $4 \text{ MB} \Rightarrow T_{\min} \approx 0.13 \text{ ms}$.
Best measured $\approx 0.94 \text{ ms}$ ($\sim 7 \times$ slower).
- **Int** (4 B): same bound, best $\approx 0.84 \text{ ms}$ ($\sim 6 \times$ slower).
- **Short** ($N = 10^4$, 2 B): $20 \text{ KB} \Rightarrow T_{\min} \approx 0.0007 \text{ ms}$.
Best measured $\approx 0.008 \text{ ms}$ ($\sim 11 \times$ slower) and dominated by timer overhead.

Interpretation:

- All measured times are well above the bandwidth bound: the kernel is not purely memory-bandwidth limited.
- Main overheads: loop control, instruction dependencies, pipeline effects, memory latency, and timer resolution for tiny runtimes.
- Loop unrolling reduces *computational* overhead but does not change the bandwidth limit.

1.6 Question 8: Why Performance Improves Then Saturates

Trend of $T(U)$:

- For small U (e.g., $1 \leq U \leq 8-16$):
 - Fewer loop iterations (N/U instead of N) \Rightarrow less branch and counter overhead.
 - More independent operations per iteration \Rightarrow better instruction-level parallelism and pipeline usage.
 - Result: strong speedup when increasing U from 1.
- For large U :
 - Loop body becomes large \Rightarrow more I-cache pressure and register pressure (spills).
 - Execution units and memory system approach saturation; further ILP does not help.
 - Small perturbations (cache effects, branch prediction, OS noise) make timings noisy; performance can degrade.

Overall, the kernel transitions from compute-bound at low U to close to a memory/latency bound at high U . The best performance is obtained with a moderate unroll factor (roughly $6 \leq U \leq 16$ depending on type and optimization level).

2 Exercise 2

2.1 Question 1: Performance at -00 vs. -02

Measured results:

- -02: $x = 131999999.662353, y = 131999999.662353$, time = 0.107 s
- -00: $x = 131999999.662353, y = 131999999.662353$, time = 0.271 s

2.2 Question 2: Analysis of Compiler Optimizations

At -00, each iteration performs two floating-point multiplications and frequent loads/stores from the stack for a , b , x , and y , with less efficient loop control.

At -02, the compiler:

- Hoists the product $a \times b$ out of the loop (loop-invariant code motion), storing it in a register.
- Uses a single accumulator in an XMM register, minimizing memory traffic.
- Holds the loop counter in a register and uses efficient branch instructions, improving instruction scheduling and instruction-level parallelism.

2.3 Question 3: Manual Optimization vs. Compiler Optimization

Manually optimized (-O0, with $a \times b$ hoisted):

- $x = 131999999.662353, y = 131999999.662353$, time = 0.265 s

Comparison:

- Original -O0: 0.271 s (slowest)
- Manual optimization: 0.265 s (improved)
- -O2: 0.107 s (fastest)

Conclusion: Manual optimization at -O0 improves performance over the naive version, but is still much slower than the compiler-optimized -O2 version. Modern compilers apply more aggressive and effective low-level optimizations than are practical to do by hand; thus, writing simple code and enabling -O2 usually yields the best performance.

3 Exercise 3

3.1 Question 1: Code Analysis

- **Strictly Sequential Parts:** The loop in `add_noise` is strictly sequential because each iteration depends on the previous one ($a[i]$ uses $a[i - 1]$). The loop in `reduction` accumulates a sum; while the provided implementation is sequential, it can be parallelized using a tree-based reduction, so it is not inherently strictly sequential in the same way as `add_noise`.
- **Parallelizable Parts:**
 - `init_b`: each $b[i]$ is independent.
 - `compute_addition`: each $c[i]$ depends only on corresponding a and b indices.
 - `reduction`: amenable to parallel reduction strategies.
- **Complexity:** All main functions (`add_noise`, `init_b`, `compute_addition`, `reduction`) have $\Theta(N)$ time complexity.

3.2 Question 2: Profiling with Callgrind

We profiled the application ($N = 10^8$) using Valgrind/Callgrind. The instruction counts (Ir) summary is:

- **Total Program:** $\approx 2.2 \times 10^9$ instructions.
- `main`: $\approx 1.1 \times 10^9$ (50%).
- `compute_addition`: $\approx 6.0 \times 10^8$ (27.27%).
- `add_noise`: $\approx 5.0 \times 10^8$ (22.73%).

The strictly sequential fraction (f_s) corresponds to `add_noise`:

$$f_s = \frac{\text{Ir}_{\text{seq}}}{\text{Ir}_{\text{total}}} = \frac{500,000,002}{2,200,140,720} \approx 0.227 \quad (22.7\%).$$

3.3 Question 3: Amdahl's Law Prediction

Using Amdahl's Law $S(p) = \frac{1}{f_s + \frac{1-f_s}{p}}$ with $f_s \approx 0.23$:

Cores (p)	Theoretical Speedup $S(p)$
1	1.00
2	1.63
4	2.38
8	3.09
16	3.63
32	3.98
64	4.18

As $p \rightarrow \infty$, the maximum speedup is bounded by $1/f_s \approx 4.35$.

3.4 Question 4: Consistency Across Problem Sizes

We repeated the profiling for $N = 5 \times 10^6$ and $N = 10^7$. In both cases, the calculated sequential fraction remained $f_s \approx 0.227$.

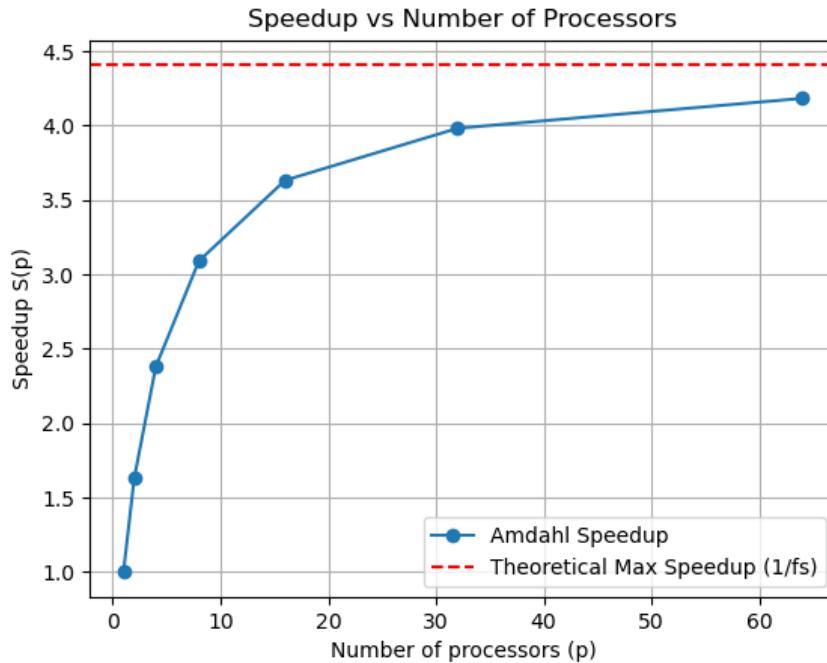


Figure 4: Theoretical Speedup vs Number of Processors (Amdahl's Law).

The speedup saturates quickly due to the fixed sequential overhead.

3.5 Question 5: Gustafson's Law

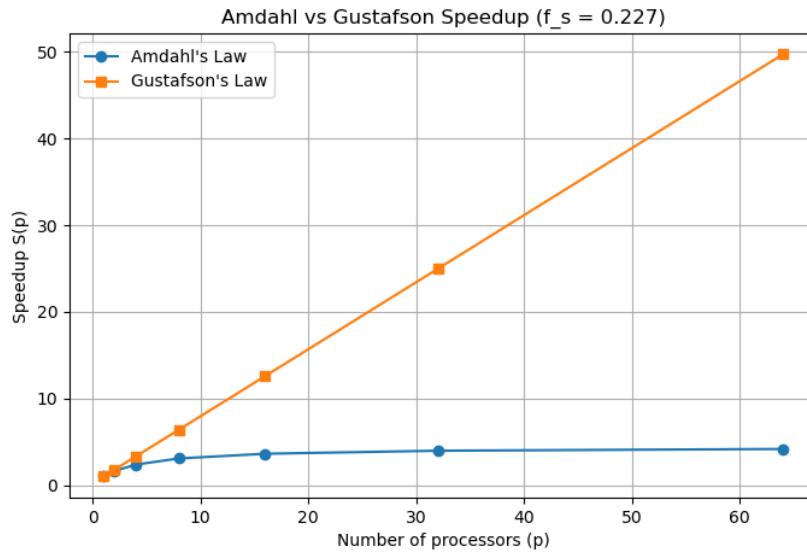
Gustafson's Law (weak scaling) assumes the problem size grows with p :

$$S_G(p) = p - f_s \times (p - 1)$$

Cores (p)	Gustafson Speedup $S_G(p)$
1	1.000
2	1.773
4	3.319
8	6.411
16	12.595
32	24.963
64	49.699

Table 1: Comparison of Scaling Laws

Law	Scaling	Growth Behavior
Amdahl	Strong	Saturates quickly (limited by f_s)
Gustafson	Weak	Grows almost linearly

**Figure 5:** Amdahl vs Gustafson Speedup.

Conclusion: Amdahl's Law shows the limits of parallelization for fixed workloads, while Gustafson's Law highlights the potential for near-linear scaling when the workload increases proportional to the processor count.

4 Exercise 4

4.1 Question 1: Profiling Analysis (N=512)

We profiled a matrix multiplication application ($N = 512$) using Valgrind/Callgrind. The instruction counts summary is:

- **Total Program:** 948,847,389 instructions.
- **matmul:** 941,888,011 (99.27%).
- **init_matrix (2 ×):** 6,815,754 (0.72%).
- **generate_noise:** 2,562 ($\approx 0.00\%$).

Within `matmul`, the innermost `k` loop accounts for approximately 99.05% of the total instructions, indicating that almost all computation time is spent in the matrix multiplication kernel.

Parallelizability Analysis:

- The outer loops over `i` and `j` in `matmul` are fully parallelizable, as each (i, j) computation is independent.
- The inner `k` loop is a reduction (sum), which is sequential per (i, j) element but can be computed in parallel for different elements.
- `init_matrix` is fully parallelizable (each element independent).
- `generate_noise` is strictly sequential due to loop-carried dependency: `noise[i] = noise[i-1] * 1.0000001`.

The sequential fraction corresponds only to `generate_noise`:

$$f_s = \frac{2,562}{948,847,389} \approx 0.00027\% \Rightarrow f_s \approx 0.003$$

This means the code is highly amenable to parallelization, with nearly all computation being parallel.

4.2 Question 2: Amdahl's Law – Strong Scaling (N=512)

Using Amdahl's Law for strong scaling with $f_s \approx 0.003$:

$$S(p) = \frac{1}{f_s + \frac{1-f_s}{p}}$$

The theoretical maximum speedup as $p \rightarrow \infty$ is $1/f_s \approx 333\times$.

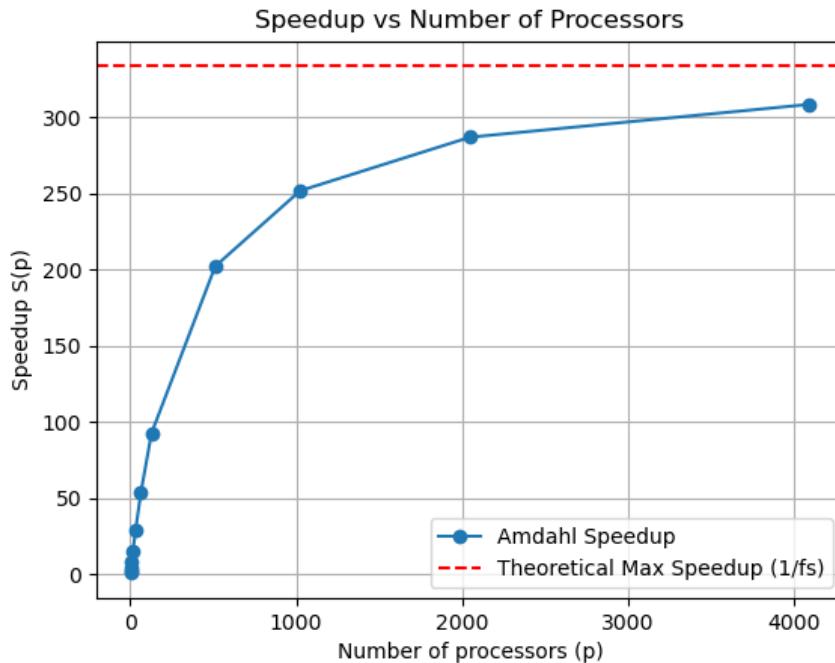


Figure 6: Amdahl's Law speedup vs. number of processors (N=512).

The code achieves near-linear speedup for a large number of processors, limited only by the small sequential fraction.

4.3 Question 3: Gustafson's Law – Weak Scaling (N=512)

Using Gustafson's Law for weak scaling:

$$S_G(p) = p - f_s \times (p - 1)$$

With $f_s \approx 0.003$, Gustafson's Law predicts nearly perfect linear scaling.

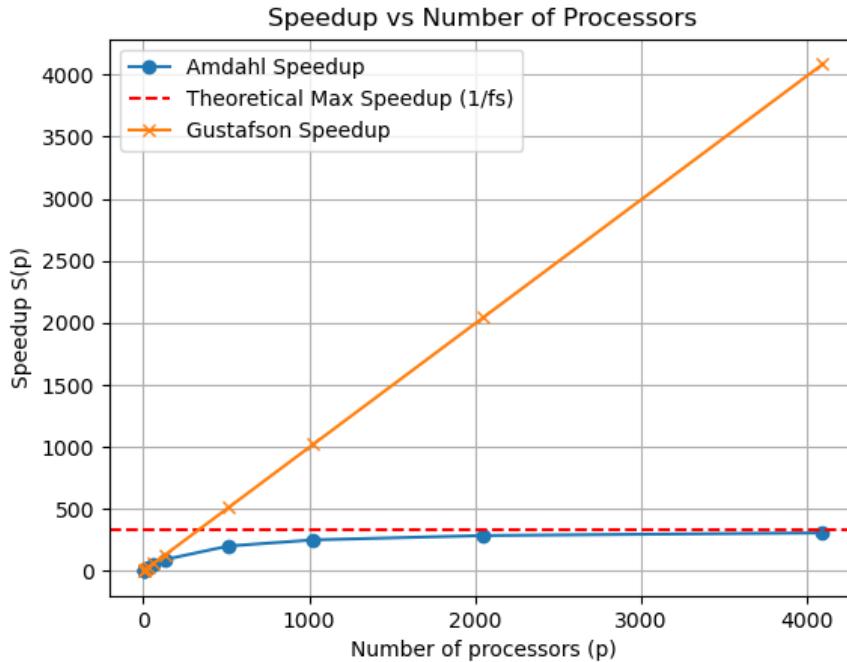


Figure 7: Amdahl vs. Gustafson speedup comparison (N=512).

4.4 Question 4: Consistency Analysis (N=128)

We repeated the profiling for $N = 128$. The instruction counts summary is:

- **Total Program:** 15,396,011 instructions.
- **matmul:** 14,828,683 (96.32%).
- **init_matrix (2 ×):** 425,994 (2.77%).
- **generate_noise:** 642 ($\approx 0.004\%$).

Sequential fraction (only generate_noise):

$$f_s = \frac{642}{15,255,319} \approx 0.000042 \quad (0.0042\%)$$

Maximum Amdahl speedup: $1/f_s \approx 23,809 \times$.

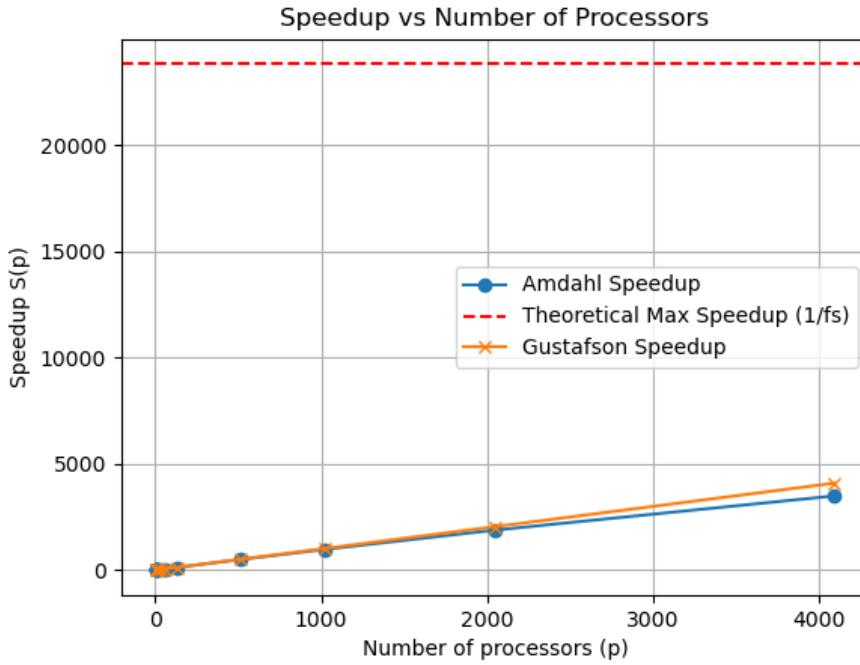


Figure 8: Amdahl vs. Gustafson speedup comparison ($N=128$).

4.5 Comparison: Exercise 3 vs. Exercise 4

Metric	Exercise 3	Exercise 4 (N=512)
Sequential function	add_noise	generate_noise
Sequential fraction f_s	22.7%	0.0003%
Parallel fraction	77.3%	99.9997%
Max Amdahl speedup ($1/f_s$)	$\approx 4.35 \times$	$\approx 333,333 \times$
Scaling at 64 cores (Amdahl)	$\approx 4.18 \times$	$\approx 64 \times$
Scaling at 64 cores (Gustafson)	$\approx 49.7 \times$	$\approx 64 \times$

Key Observations:

- Sequential fraction impact:** Exercise 3 has a significant sequential portion (22.7%), severely limiting speedup under Amdahl's Law. Exercise 4's sequential fraction is negligible (< 0.001%), allowing near-linear scaling.
- Complexity differences:** In Exercise 3, all functions scale as $O(N)$, so the sequential fraction remains constant regardless of problem size. In Exercise 4, `matmul` scales as $O(N^3)$ while `generate_noise` scales as $O(N)$, causing the sequential fraction to decrease as N grows.
- Amdahl vs. Gustafson:** For Exercise 3, Gustafson's Law provides much better scaling predictions than Amdahl's Law due to the large f_s . For Exercise 4, both laws predict nearly identical (linear) scaling because $f_s \approx 0$.
- Practical implications:** Exercise 3 represents a poorly parallelizable workload where weak scaling is essential for good performance. Exercise 4 represents an ideal parallelizable workload where strong scaling is achievable.

Conclusion: The sequential fraction f_s is the critical factor determining parallel scalability. Algorithms with $O(N^k)$ parallel work and $O(N)$ sequential overhead (like matrix multiplication with $k = 3$) become increasingly parallelizable as problem size grows, while algorithms where all components scale equally (like Exercise 3) maintain a fixed scalability limit.